

Can Social Proof Fool You?

How Engagement Metrics Can Affect AI vs Human Content Detection

Presenter: Vishakha Agrawal

Introduction

RM

Ranjani Mani

Director and Country Head, Gen...

8h · 🌐

Follow

...

In school, I used to cover my ears and rock back and forth to block out noise and study — peak nerd energy.

My friends still tease me about it.

Last weekend I met old school friends — the kind who knew us before job titles, metrics, and calendar chaos.

There's something uncanny about meeting people you've known since you were three: decades pass, and within minutes the rhythm returns.

Time folds. The laughter is effortless — and so is the mirror it holds up: who you were, who you became, and what you may have lost along the way.

We've all taken different paths — some brave, some accidental, all real. That contrast is part of the beauty.

Maybe growth isn't becoming someone new, but remembering your roots — choosing which parts of yourself to keep as you evolve.

The people who knew our unpolished selves still see us most clearly. Who would your 15-year-old self be proud to meet today?

Ranjani Mani

#reviewswithranjani

#Technology | #Books | #BeingBetter

RM

Ranjani Mani

Director and Country Head, Gen...

8h · 🌐

Follow

...

I did this thing in school

Cover my ears and rock to cut out the noise around to focus on my books, the nerd that I was

My friends still tease me about it

Met a few school friends last weekend — people who knew other before the titles, before the metrics, before life started needing calendars.

There is something strange about meeting people you know from when you were 3 years old - you meet after decades and yet fall right into rhythm in minutes

It's strange how time folds when you meet them. The laughter returns effortlessly, but so does a mirror — showing who you were, who you became, and what you may have lost in between.

We've all taken different routes since then — some brave, some accidental, all real. And that contrast is what makes it beautiful.

Maybe growth isn't about becoming someone new. It's about remembering your roots — and choosing which parts of yourself to keep as you evolve.

Funny how the people who knew our unpolished selves still see us the dearest. Who would your 15-year-old self be proud to meet today?

Ranjani Mani

#reviewswithranjani

#Technology | #Books | #BeingBetter

Can We Tell the Difference?

- Are we any good at it? And what influences our guesses?
- AI-generated text is becoming indistinguishable from human writing. This creates new challenges for trust, misinformation, and authenticity online.
- We don't make these judgments in a vacuum. We're influenced by design, the author, and especially, social engagement cues.

Literature Review

The Detection Problem

Humans Are Unreliable Detectors (Jakesch et al., 2023) [Link](#)

- **What they found:** People are fundamentally bad at distinguishing AI from human self-presentations, with accuracy hovering around 50% (a coin flip).
- **Why:** We rely on "flawed heuristics" (e.g., "AI is formal," "humans use 'I'"). These instincts are wrong, and AI can even be optimized to seem "more human than human."

Literature Review

AI Detectors Are Also Unreliable (Weber-Wulff et al., 2023) [Link](#)

- **What they found:** A comprehensive test of 14 tools (including Turnitin) found they are "neither accurate nor reliable."
- **Why:** They are easily fooled by simple "obfuscation techniques" like using a paraphrasing tool (Quillbot) or even machine translation (which creates false positives for non-native English speakers).

Literature Review

The Heuristic Solution

We Use Likes as a Shortcut (Metzger et al., 2010) [Link](#)

- **What they found:** In an information-abundant environment (the internet), people don't have time to systematically check facts.
- **Why:** We use cognitive heuristics to judge credibility. A primary shortcut is the **Endorsement Heuristic**—we trust what is popular.

Literature Review

Likes Are Proven to Bias Our Judgment (Muchnik et al., 2013) [Link](#)

- **What they found:** In a large-scale experiment, a *single, artificial* initial up-vote on a comment created a "positive herding" cascade.
- **Why:** This social cue caused the comment's final score to be **25% higher** than an identical (control) comment. This proves that "likes" actively bias our perception of quality.

Literature Review

Labels Create *Anti-AI* Bias (Zhu et al., 2025) [Link](#)

- **What they found:** In a blind test, people couldn't tell AI from human text. But when *labels* were applied, they overwhelmingly favored "Human Generated" text by over 30%, **even when the labels were deliberately swapped.**

Literature Review

- Liu, M., & Conrad, F. G. (2016). An experiment testing six formats of 101-point rating scales. *Computers in Human Behavior*, 55, 364–371. [Link](#)
- Boot, A. B., Dijkstra, K., & Zwaan, R. A. (2021). The processing and evaluation of news content on social media is influenced by peer-user commentary. *Humanities and Social Sciences Communications*, 8(1), 209. [Link](#)
- Fleming, S. M. (2023). Metacognition and confidence: A review and synthesis. *Annual Review of Psychology*. [Link](#)

Research Question

Does social media engagement (likes) affect people's ability to distinguish AI-generated from human-written content?

Hypotheses

H1: Engagement Effect on Accuracy

Does high popularity (likes) reduce detection accuracy?

H2: Error Asymmetry

Does engagement change the *type* of errors we make?

More AI→Human errors than Human→AI errors?

H3: Confidence–Accuracy Relationship

How does confidence relate to correctness—does high popularity inflate confidence?

Experiment Design

Design

A 2x2 Within-Subjects Experiment

Independent Variables (IVs):

- IV1: Text Version: Original (Human) vs Rephrased (AI) [Nominal]
- IV2: Engagement Level: Low (3–50 likes) vs High (6,000–50,000 likes) [Ordinal]

Dependent Variables (DVs):

- DV1: Attribution Accuracy: Correct/Incorrect [Nominal/ Binary]
- DV2: Confidence Rating: 0–100 scale (steps of 10) [Interval]

Design

Additional Measures:

- Dwell Time: Stimulus viewing duration in ms [Ratio]
- Response Time: Decision time in ms [Ratio]

Operationalization

- Human content: Real LinkedIn posts by verified authors
- AI content: Claude 4 Sonnet rephrased for readability (same semantic meaning, similar length)
- Likes: Sampled log-uniformly within ranges, displayed prominently

Conditions and Sampling

Conditions per participant (20 trials total):


- 5× Human–Low, 5× Human–High, 5× AI–Low, 5× AI–High

To prevent order effects (like fatigue or practice), the presentation of the four conditions was randomized for each participant.

Sampling:

- N = 35 student participants (Convenience Sampling – Recruited from Workspace, THub, and Parijat)
- Gender: 14 Males, 21 Females
- Age Group: 18–23 years

Trial Flow



Consent

You will view social media posts and judge whether they were written by an AI (Artificial Intelligence) or a human, then rate your confidence. Your responses are anonymous.

Do you consent to participate?

☒ Yes, I consent

Continue

Participant ID: c568bdb4-7104-4686-9b3c-a5c0e4c02d3b

Instructions

- Read each post carefully.
- When you are ready, click on the Answer button below each post
- Answer whether it was written by **AI** or a **Human**.
- Rate your confidence from **0** to **100** in steps of **10**.

Start

Participant ID: c568bdb4-7104-4686-9b3c-a5c0e4c02d3b

DB

Darius Burschka



3rd +

Professor CIT (TUM) - Visual Analysis of Dynamic Scenes

1d · Edited ·

Follow



Mixed feelings. Great to see a **#novel** method tackle an old problem — but here's my first concern. Why is everything tossed into the same “**#AI**” bucket, as if a **#LLM** was always involved? We should clearly distinguish learning approaches; otherwise useful applications may keep weak ones (popular, but unhelpful) alive.

THOR AI claims to avoid the classic bottlenecks — where traditional integration would take longer than the age of the universe — by using *tensor train cross interpolation*. In short, it decomposes a high-dimensional equation into smaller, linked tensors, turning an intractable computation into a manageable one.

My question: does this method actually require AI, or could a classical system run it just as well? It reads like a novel technique repackaged as “AI hype” for funding.

<https://lnkd.in/ecSxxJk4>

24,946

Answer

Your judgment

Trial 1 / 20

Who do you think wrote this post?

☐ AI ☒ Human

How Confident are you in your response? (0–100)

0 = least confidence, 100 = most confidence. Steps of 10.



Your value: **90**

Next

Experiment 1: Human vs. AI Performance on Text Classification															
Participant Information				Task Details				Performance Metrics				System Configuration			
A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P
participant_id	trial_index	post_id	source_true	text_version	likes_level	likes_value	order	response_choice	confidence_0_100	dwell_ms	rt_ms	started_at	device	viewport_w	viewport_h
3a03ad1d-de94-49	0	stefano_rp_low	AI	Rephrased	Low	2	0	Human	80	55255	11815	2025-11-07T10:30:00	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	1	travel_rp_high	AI	Rephrased	High	1762	1	Human	70	62591	15231	2025-11-07T10:30:01	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	2	live_shopping_og_low	Human	Original	Low	12	2	Human	80	49580	7672	2025-11-07T10:30:02	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	3	madhav_og_high	Human	Original	High	2735	3	AI	70	46384	5168	2025-11-07T10:30:03	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	4	evolvingai_rp_high	AI	Rephrased	High	2773	4	AI	80	11736	7064	2025-11-07T10:30:04	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	5	aman_og_low	Human	Original	Low	1	5	AI	90	8617	6800	2025-11-07T10:30:05	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	6	nicholas_og_high	Human	Original	High	20392	6	AI	80	42807	4801	2025-11-07T10:30:06	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	7	barnacles_rp_low	AI	Rephrased	Low	0	7	Human	90	19650	4728	2025-11-07T10:30:07	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	8	garius_rp_low	AI	Rephrased	Low	1	8	Human	70	28338	4672	2025-11-07T10:30:08	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	9	jaineel_og_high	Human	Original	High	1339	9	Human	70	33401	3856	2025-11-07T10:30:09	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	10	russelladams_og_low	Human	Original	Low	1	10	AI	60	20873	9153	2025-11-07T10:30:10	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	11	amitgupta_rp_high	AI	Rephrased	High	4886	11	AI	70	24265	5151	2025-11-07T10:30:11	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	12	reema_og_high	Human	Original	High	29692	12	AI	90	25602	5658	2025-11-07T10:30:12	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	13	sunjita_rp_high	AI	Rephrased	High	1792	13	Human	60	30120	5705	2025-11-07T10:30:13	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	14	sudhya_rp_low	AI	Rephrased	Low	0	14	Human	80	33264	3672	2025-11-07T10:30:14	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	15	ranjani_og_low	Human	Original	Low	0	15	AI	70	23040	4792	2025-11-07T10:30:15	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	16	amit_rp_high	AI	Rephrased	High	8412	16	Human	80	39455	4552	2025-11-07T10:30:16	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	17	paulo_og_low	Human	Original	Low	9	17	AI	70	16272	4064	2025-11-07T10:30:17	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	18	gary_rp_low	AI	Rephrased	Low	1	18	Human	70	25967	4673	2025-11-07T10:30:18	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671
3a03ad1d-de94-49	19	charles_og_high	Human	Original	High	4356	19	AI	70	34816	9816	2025-11-07T10:30:19	Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36	1536	671

Ethics

Informed Consent

Participants informed about study purpose, duration (~20 min), and voluntary participation

Anonymity & Privacy

Random IDs only (no names, emails, or personal data)

- No deception
- Right to withdraw anytime

Minimal Risk

Task stimuli similar to everyday social media browsing; no sensitive content

Sources of Variability & Control

Random Variables

- Individual differences in reading speed, attention, fatigue
 - Random fluctuations/ interpretation in confidence ratings
 - Testing environment: Participants completed study at different times of day (morning, afternoon, evening), in varied settings, with different levels of distraction, doing different activities.
-
- Control Strategy: Larger sample size, same setting

Sources of Variability & Control

Systematic/Constant Variables

- Post content quality: Some posts may be inherently easier/harder to classify.
- Control Strategy: Within-subjects design (each person is their own control) and randomization.

Sources of Variability & Control

Potential Confounds:

- AI familiarity: Tech-savvy participants may be better at detection (not measured)
- Platform familiarity: LinkedIn users may have different detection strategies (not assessed)

Sources of Variability & Control

Potential Confounds We Controlled

Confound	How It Could Bias Results	Our Control
Post topic	Some topics easier to detect AI	Randomized assignment of posts to conditions
Text length	Longer posts = more time to judge	Original vs Rephrased matched for semantic content
Participant motivation	Declining engagement over time	Short experiment (~15 min)

Validity

1. Internal Validity

Strengths:

Within-subjects design controls for individual differences

Random trial order prevents order effects

Threats: Only 20 trials, minimal observation and learning

Validity

2. External Validity

Strengths:

Real LinkedIn posts (authentic content)

Realistic engagement metrics

Limitations:

Low ecological validity

N=35 is too small to generalize to all social media users.

Convenience sample (does not represent all demographics and general population in terms of age and type)

Lab setting $\neq$ natural social media scrolling where different cues mediate interaction

Only LinkedIn posts (not Twitter, Facebook, etc.)

Validity

3. Construct Validity

Accuracy: Binary choice (AI/Human) directly measures attribution

Engagement: Like counts clearly operationalize social engagement

Confidence: 0–100 scale for metacognitive judgments

4. Statistical Conclusion Validity

Strengths:

- Appropriate tests (RM-ANOVA for within-subjects)

Limitations:

- Small sample (N=35)

- Normality violations (but robust with this N)

Reliability

1. Internal Consistency

Do multiple trials within the same condition measure consistently?

Design: 5 trials per condition per participant

Result: Multiple observations allow stable accuracy estimates per condition

Trade-off: More trials = better reliability, but longer experiment (fatigue risk)

2. Test-Retest Reliability

Would participants get similar results if tested again?

Not assessed in this study (single session, no follow-up)

Future: Could retest subset of participants after 1 week to check stability

Reliability

3. Inter-Rater Reliability

Not applicable – objective binary outcome (correct/incorrect based on true source)
No subjective coding or human judgment required

4. Measurement Reliability

Are our instruments consistent?

Accuracy: Objective (response matches ground truth) → Perfect reliability

Confidence: Self-report scale

Timing: Automated (millisecond precision) → Highly reliable

Results Overview

Descriptive Statistics

Sample Characteristics

Final N = 35 participants (after deduplication & quality control)

Total trials = 700 (35 participants × 20 trials)

100% completion rate (no missing data)

Device: All desktop/laptop (Chrome browser)

Median viewport: 1536×671 pixels


Overall Performance

Measure	Result
Overall Accuracy	54.43% (barely above chance)
Response Distribution	50.7% "Human", 49.3% "AI"
Mean Confidence	69.4 / 100
Mean Dwell Time	29.6 seconds
Mean Response Time	5.7 seconds

Accuracy by Condition


Condition	Accuracy
Original-Low	57.7%
Original-High	52.6%
Rephrased-Low	49.1%
Rephrased-High	58.3%

1. Accuracy by Condition



Error Analysis

- AI→Human errors (False Negatives)
46.3% (162/350)
- Human→AI errors (False Positives)
44.9% (157/350)
- Difference: 1.4% (nearly balanced)



Inferential Statistics

H1 (Engagement Effect)

H1: Does high engagement (likes) reduce accuracy?

Test: 2×2 Repeated Measures ANOVA

- DV: Accuracy (proportion correct per condition)
- IVs: Text Version (Original vs Rephrased) × Engagement (Low vs High)

Results

Effect	F-statistic	p-value	η^2_G	Interpretation
Text Version	F(1,34)=0.089	p=0.767	0.001	No effect [ns]
Engagement	F(1,34)=0.265	p=0.610	0.002	No effect [ns]
Version × Engagement	F(1,34)=3.049	p=0.090	0.021	Marginal (Not significant)


Interpretation

Main Effects: No significant effect of text version or engagement alone

Interaction ($p=0.090$): Marginal crossover, no significant interaction

H1 NOT SUPPORTED: High engagement doesn't reduce accuracy overall, but may affect human vs AI detection differently (needs replication with larger sample)

Interaction: Text Version × Engagement (p = 0.090)



H2: Are errors asymmetric?

Do people make more AI→Human errors (false trust) than Human→AI errors (false suspicion)?

Test: Paired t-test

- Compared error counts within participants
- AI errors (missed detecting AI) vs Human errors (falsely called Human "AI")

Results

Error Type	Mean Errors (out of 10)	Rate
AI→Human (False Negatives)	4.63	46.3%
Human→AI (False Positives)	4.49	44.9%
Difference	0.14	1.4%

Interpretation

$t(34) = 0.299$, $p = 0.767$, Cohen's $d = 0.050$


Essentially no difference between groups

No significant asymmetry

H2 NOT SUPPORTED: Errors are balanced

- People are equally bad at detecting AI content as they are at falsely labeling human content as AI
- No systematic bias in error patterns
- The 1.4% difference is not statistically or practically significant

H2: Error Asymmetry per Participant (Points above diagonal = more AI→Human errors)



H3: Does confidence relate to accuracy?

Test 1: Paired t-test (Confidence by Correctness)

Compared confidence ratings when correct vs incorrect (within participants)


	Confidence (0-100)
When Correct	69.2
When Incorrect	69.8
Difference	-0.6 (not significant)

Statistical Test:

$t(34) = -0.236, p = 0.815$

Interpretation: Participants equally confident whether right or wrong.

5. Confidence by Correctness



Test 2: Calibration

Brier Score = 0.307

Measures overall prediction accuracy (0 = perfect, 0.25 = chance)

0.307 = worse than chance → confidence doesn't predict correctness

Mean Calibration Error = +0.15


Confidence (69.4%) - Accuracy (54.4%) = +15%

Overconfident by 15 percentage points

Positive = overconfident | Negative = underconfident | Zero = perfect

Participants feel confident but perform poorly.

Calibration



Test 3

2×2 RM-ANOVA (Confidence ~ Correctness × Engagement)

Effect	F-statistic	p-value	Interpretation
Correctness	F(1,34)=0.043	p=0.837	No effect [ns]
Engagement	F(1,34)=0.010	p=0.922	No effect [ns]
Correctness × Engagement	F(1,34)=0.338	p=0.565	No effect [ns]

No significant effects – Confidence is unaffected by correctness or engagement

Interpretation

H3 NOT SUPPORTED:

- Confidence remains stable (~69–70) regardless of actual performance
- Overconfident by 15% on average
- Engagement doesn't inflate or deflate confidence

Exploratory Analysis - Signal Detection Theory

Separates discrimination ability from response bias


Results

Condition	d' (sensitivity)	c (bias)
Overall	0.223	0.018
Low Likes	0.173	0.108
High Likes	0.274	-0.072

Interpretation

Very weak discrimination ($d' = 0.22$, close to chance)
Neutral response bias ($c \approx 0$, no systematic preference)
No engagement effect on either sensitivity or bias

SDT: Hit Rate vs False Alarm Rate



Qualitative Insights - Detection Strategies

Post-experiment open-ended informal responses: "What cues did you use to judge if content was AI or Human?"

Emergent Themes

Surface-Level Features

- Emojis: Frequency and placement patterns
- Punctuation: Excessive ellipses (...), quotation marks, exclamation points
- Symbols & Tags: Unnecessary hashtags, special characters
- Line spacing: Consistency and formatting patterns

Qualitative Insights - Detection Strategies

Linguistic Patterns

- Repetition & Redundancy: Restating the same point multiple times
- Structural consistency: Uniform sentence length or paragraph structure
- Complexity: "Too perfect" grammar vs natural errors

Content Structure

- Explanation patterns: Stating something, then over-explaining it
- Formulaic writing: Predictable intro-body-conclusion structure

Discussion and Limitations

Key Findings

- Participants were unable to reliably detect AI content (~54% accuracy)
- High engagement does not reduce accuracy overall
- **No significant interaction:** Likes may affect human vs AI detection differently
- No systematic error bias in either direction
- Confidence unrelated to correctness

Theoretical Implications

- Social proof effects are more complex than expected
- AI detection could be challenging, regardless of context cues

Practical Implications

- Users cannot self-detect AI content reliably → need explicit labeling/disclosure

Misinformation & Content Moderation

- Relying on user reports for AI content may be ineffective (poor detection)
- Automated detection systems remain essential
- AI literacy programs can teach which cues are reliable vs misleading

Limitations - Sampling

Small Sample Size (N = 35)

- Underpowered for detecting small effects
- Very low generalizability

Convenience Sampling

- Students and participants familiar with technology
- Does not represent general population (age, education, AI familiarity)

Limitations - Design

Stimuli Constraints

- Only LinkedIn posts (professional content) → limited generalizability to other platforms
- Only one AI model (Claude 4 Sonnet) → different models may produce more/less detectable text
- 20 base posts → limited content diversity

Engagement Manipulation

- Participants may not attend to like counts as they would during natural browsing
- Lab setting lacks ecological validity of real social media context

Limitations - Validity

Ecological Validity

- Single-session experiment $\neq$ longitudinal exposure to AI content
- Static posts $\neq$ dynamic social media feeds with multiple cues

Future Directions

- Larger sample with diverse demographics
- Multiple content types (Twitter, articles, emails)
- Eye-tracking to verify attention to engagement metrics
- Longitudinal design (do people learn over time with repeated exposure?)
- Individual Differences – Do AI familiarity, education, or domain expertise affect detection?
- Systematic thematic analysis and coding of detection strategies
- Test interventions (e.g., AI literacy training)
- Field studies in real social media environments
- Test detection in high-stakes contexts (news, health information)

Thank You!

Do you have any questions?

Resources: [Link to Analysis Notebook](#)